

AI Industry Highlights

August 26, 2026 | Daily Briefing

Hardware & Chips

Apple M6 / M5 Ultra	M6: 2nm process, 12-core CPU + 12-core GPU. M5 Ultra: up to 80-core GPU, 512GB unified memory, 4.5x AI compute of M3 Ultra.
OpenAI Jalapeño Chip	First custom inference chip taped out with Broadcom on TSMC N3P. 13.4 PFLOPs MXFP4 @ 700W, 1.5-1.9x perf-per-watt vs Nvidia Rubin.
Nvidia Groq 3 LPX	Dedicated inference accelerator enters mass production. 3,400 tok/s on Gemma 4 31B @ 100K context. Nebius is first cloud customer.

Models & Platforms

Qwen 3.8-Flash-Next	Alibaba's next-gen open MoE model (~125B total, ~6B active/token) previewing Qwen 4 architecture. Community prepping day-zero support.
Anthropic Unified Memory	Claude chat and Cowork now share memory across surfaces by default. Sensitive categories blocked; enterprise admins can disable.
Ox Alpha (GLM-5.3)	Anonymous model on OpenRouter with 1M-token context. Scores 80% on DeepSWE benchmark, ahead of Claude Fable 5 (65%) and GPT-5.6 (52%).

Industry & Startups

SpaceX acquires Anysphere	\$600B deal for Cursor parent company -- largest tech acquisition in history. Cursor annualized revenue exceeds \$4B.
Stability AI	Raised \$76M from EA, Sony Music, UMG, WMG, and AMD for creative-production products and professional services.
Gatik	Autonomous trucking startup raises \$200M, now operating 41 driverless box trucks for PepsiCo across Dallas, Phoenix, and Arkansas.

HARDWARE	LLM	ROBOTICS	CHIPS	STARTUPS	OPEN SOURCE
----------	-----	----------	-------	----------	-------------

Sources: AI Weekly, Sina AI, BBC News, Reuters, TechCrunch | Generated for briefing purposes only